

vclsi-03-Mueller-Lordick-Aborageh

May 9, 2019

Simon Müller s6simue@uni-bonn.de
Tomas Lordick s0thlord@uni-bonn.de
Mohamed Aborageh mohamed.salah1@aiesec.net

0.1 Exercise 1

a)

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
from sklearn.decomposition import PCA
from sklearn.preprocessing import scale
from sklearn.preprocessing import StandardScaler
import seaborn as sb
import numpy as np

In [3]: df = pd.read_excel("breast-cancer-wisconsin.xlsx") #read the data set
df=df.interpolate(method='nearest') #interpolate missing values

In [4]: df=df.drop("code", axis=1) #drop the code column
df.head()
features = df.loc[:, list(df)[-1]].values #split the data into features set and class
klasse = df.loc[:, ["class"]].values

In [5]: features_scaled = StandardScaler().fit_transform(features) # scaling the features set

In [6]: pca = PCA() # doing the actual pca
principalComponents = pca.fit_transform(features_scaled)
principalDf = pd.DataFrame(data = principalComponents
                           , columns = [i for i in range(1,10)])

In [7]: pca.fit(features_scaled) # fit teh data to the pca

Out[7]: PCA(copy=True, iterated_power='auto', n_components=None, random_state=None,
          svd_solver='auto', tol=0.0, whiten=False)

In [8]: # plot the cumulated variance.
variance =pca.explained_variance_ratio_ # calculate variance ratios
var=np.cumsum(np.round(pca.explained_variance_ratio_, decimals=3)*100) #rounded to per
```


```

plt.ylabel('% Variance Explained')
plt.xlabel('n +1 components ')
plt.xticks([i for i in range(10)])
plt.title('PCA Analysis')
plt.ylim(30,100.5)
print (var)
plt.plot(var)

```

```
[ 65.4  74.   80.   85.2  89.4  92.8  96.1  99.  100. ]
```

Out[8]: [



In [9]: # the first 6 components are enough to cover >= 90 %.

b)

In [11]: finalDf = pd.concat([principalDf, df[['class']]], axis = 1) #concatenate scaled features

In [12]: finalDf.head()

```

Out[12]:
   1      2      3      4      5      6      7  \
0 -1.458989 -0.103530 -0.577553 -0.023162 -0.151570  0.146176  0.300721

```

1	1.457655	-0.553216	0.310822	-0.655750	1.691574	-0.561744	-0.924616
2	-1.582918	-0.073038	0.040621	-0.100203	-0.067063	-0.211299	0.273373
3	1.501772	-0.545678	-0.649455	1.431360	-0.447196	-0.159949	-0.236603
4	-1.333143	-0.085412	0.028890	-0.300618	-0.151717	0.515047	0.160909

	8	9	class
0	0.435618	-0.001557	2
1	0.352840	0.012818	2
2	0.233932	0.015796	2
3	-1.593375	0.186781	2
4	0.435507	-0.034480	2

```
In [13]: first_five=finalDf.drop(columns=[6,7,8,9]) # drop the components 6,7,8,9 to keep the .
```

```
In [14]: first_five.head()
```

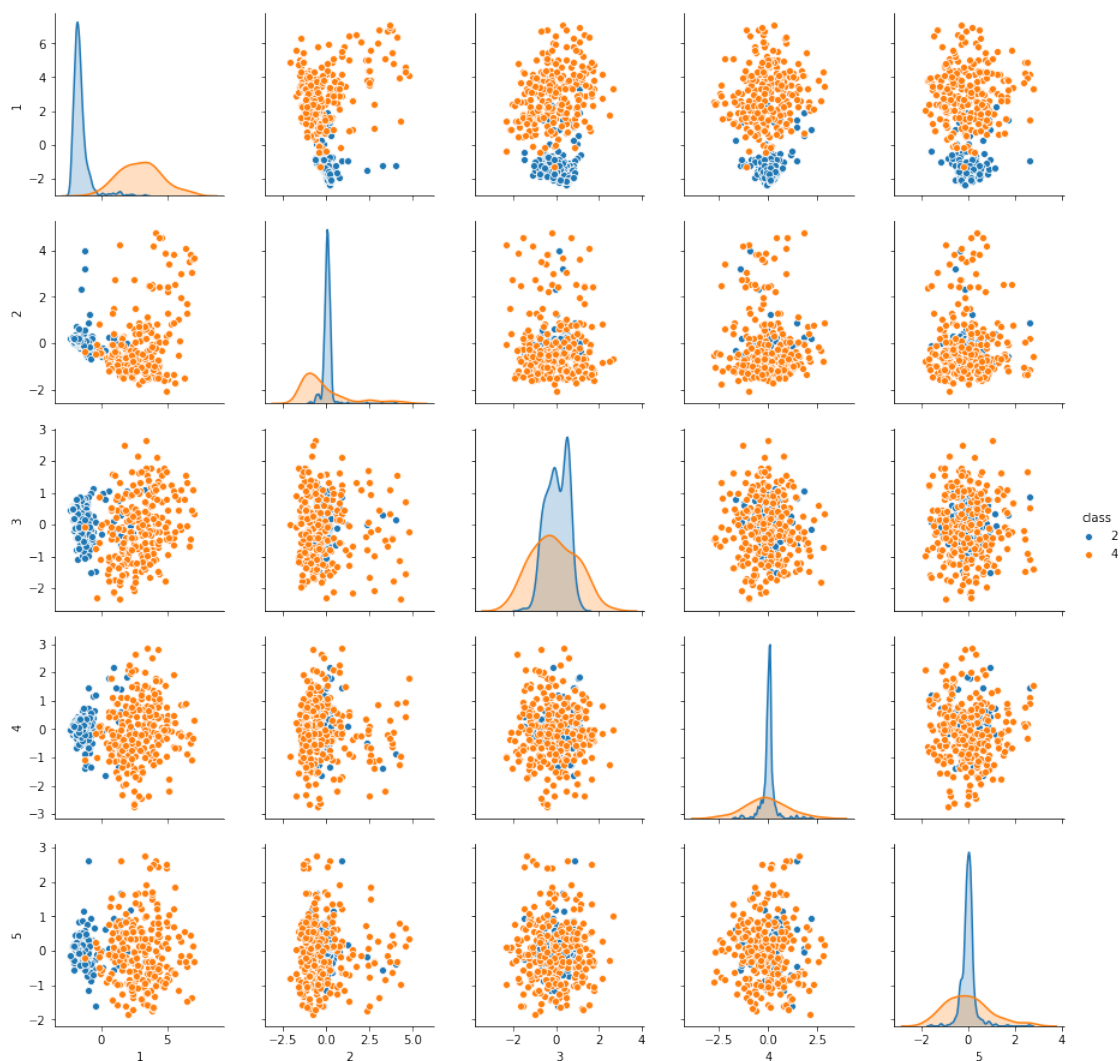
```
Out[14]:
```

	1	2	3	4	5	class
0	-1.458989	-0.103530	-0.577553	-0.023162	-0.151570	2
1	1.457655	-0.553216	0.310822	-0.655750	1.691574	2
2	-1.582918	-0.073038	0.040621	-0.100203	-0.067063	2
3	1.501772	-0.545678	-0.649455	1.431360	-0.447196	2
4	-1.333143	-0.085412	0.028890	-0.300618	-0.151717	2

```
In [15]: sb.pairplot(first_five,vars=[1,2,3,4,5], hue="class") #plot is using seaborn
```

```
F:\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-tuple seq
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

```
Out[15]: <seaborn.axisgrid.PairGrid at 0x20652223390>
```



```
In [16]: print (pca.explained_variance_)
          print (pd.DataFrame(pca.components_,columns=list(df)[:5]))
```

```
[5.89047615  0.77744685  0.53964716  0.46827076  0.38087324  0.30774761
  0.29723871  0.26212251  0.089071   ]
   thickness  uniCelS  uniCelShape  marAdh  epiCelSize  bareNuc  \
0    0.302844  0.381430    0.377742  0.333094    0.336632  0.331780
1   -0.142161 -0.047669   -0.083087 -0.046233    0.164846 -0.256308
2   -0.862991  0.012783   -0.040017  0.427623    0.103621  0.017284
3   -0.106419  0.206584    0.178597 -0.449796    0.372007 -0.554646
4   -0.069211  0.136537    0.103516 -0.035302    0.682886  0.128276
5    0.254525  0.142858    0.066324  0.655310   -0.085412 -0.641377
6   -0.065129  0.186261    0.133115 -0.232489   -0.174851 -0.267513
7    0.241196 -0.439179   -0.584973  0.122137    0.448963 -0.109863
8    0.002544  0.736734   -0.664741 -0.045680   -0.067884  0.069606
```

	blaChroma	normNuc	mitoses
0	0.346282	0.336174	0.229855
1	-0.229701	0.025845	0.906741
2	0.194083	0.121386	-0.084739
3	0.000763	0.451138	-0.249472
4	-0.255242	-0.632166	-0.130518
5	-0.155318	-0.158974	-0.113679
6	0.726149	-0.485984	0.153899
7	0.411936	0.081437	-0.043292
8	-0.058328	0.018414	-0.007395

```
In [17]: # I am not sure what PCA "mode" means, but this would be my guess:
# The first component covers ~ 65 % of the variance. Its clearly visible in the first
# Regarding this component, we see that:
# uniCelS shows the highest weight in the first principal component (~ 0.38)
# thickness reveals the lowest weight in the first principal component (0.30 )
```

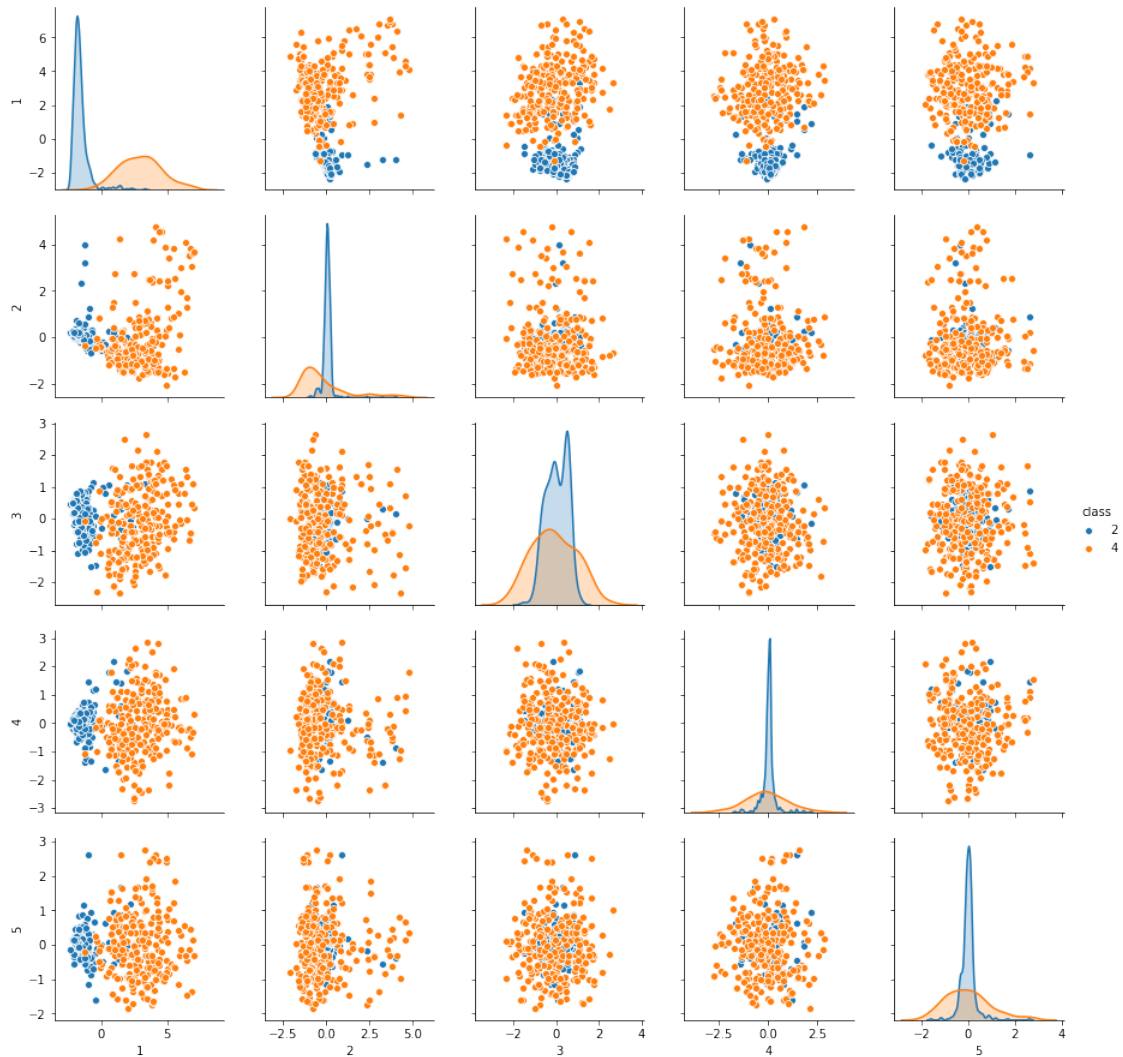
d)

```
In [18]: #outlier removal: row with index 1.
#Will probably not make any major difference. To archieve that, more outlier samples .
#Just to demonstrate:
first_five_removed=first_five.drop([1])
```

```
In [19]: #replot
sb.pairplot(first_five,vars=[1,2,3,4,5], hue="class")
```

```
F:\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-tuple seq
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

```
Out[19]: <seaborn.axisgrid.PairGrid at 0x2064fadbc50>
```



e)

Pre-processing is always necessary even for this dataset - that's why I've already made it as you can see in a). We're doing that in order to remove mathematically the sources of unwanted variations (QUOTED). If the relationships between the variables analyzed are non-linear, the values of correlation coefficients can be lower. So, it is sometimes useful to transform the original variables prior to the Principal Component Analysis to "linearize" these relationships.

0.2 Exercise 2

a)

- if you set the perplexity to 30, you have the same number of points which results in a uniform distribution that does not yield good results. This does not happen when perplexity $< n$, which explains the much better results

- t-SNE cares more about close neighbors than more distant data points. At the edges and corners you have more neighbors which pushes these points closer together
- a neighborhood of only two points is not enough to hold a square structure together
- perplexity has to be bigger than three. You have to look at enough neighbors around a point to bridge the gaps between the different circle elements. This does only happen if also look at points outside the local cluster.

b)

```
In [20]: import numpy as np
import matplotlib
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.decomposition import PCA
from sklearn.manifold import Isomap, TSNE


In [21]: df = pd.read_excel('Data_Cortex_Nuclear.xls')
df = df.interpolate(limit_direction='backward')
sc_1 = df.loc[df.loc[:, "class"] == "t-CS-s", :]
sc_2 = df.loc[df.loc[:, "class"] == "c-CS-s", :]
df = pd.concat([sc_1, sc_2])
df_num = df.select_dtypes(['number'])
df_num.head()

print('Data len: ', len(df_num))
```

Data len: 240


```
In [22]: #PCA
pca = PCA(n_components=2)
df_red = pd.DataFrame(
    data=pca.fit_transform(df_num),
    columns=['A', 'B'],
    index=df_num.index
)
df_red['class'] = df['class']
sns.scatterplot(data=df_red, x='A', y='B', hue='class')
```

Out[22]: <matplotlib.axes._subplots.AxesSubplot at 0x2065585fd30>



```
In [23]: #isomap
iso = Isomap(10)
df_red = pd.DataFrame(
    data=iso.fit_transform(df_num),
    columns=['A', 'B'],
    index=df_num.index
)
df_red['class'] = df['class']
sns.scatterplot(data=df_red, x='A', y='B', hue='class')
```

```
Out[23]: <matplotlib.axes._subplots.AxesSubplot at 0x20655a74198>
```

I would choose the isomap visualization because it stretches the data better along the x-axis which makes it more readable. Also it is easier to spot certain clusters.

c)










```
In [24]: data = pd.read_excel('breast-cancer-wisconsin.xlsx')
data = data.drop('code', axis=1)
data = data.interpolate(limit_direction='backward')
data_num = data.select_dtypes(['number'])
data_label = data['class']

inits = ['random', 'pca']
perplexity_params = [5, 10, 20, 30, 40]

fig, axs = plt.subplots(nrows=10, figsize=(15,100))
count = 0
for ini in inits:
    for perp in perplexity_params:
        print(f'Generating plot with perplexity "{perp}" and init-mode "{ini}"')
        tsne = TSNE(perplexity=perp, init=ini)
        data_red = pd.DataFrame(
            data=tsne.fit_transform(data_num),
            columns=['A', 'B'],
            index=data_num.index
        )
```

```
data_red['class'] = data['class']
sns.scatterplot(data=data_red, x='A', y='B', hue='class', ax=axes[count])
axes[count].set_title(f'Plot with perplexity {perp} and init-mode {ini}')
count += 1
```

```
Generating plot with perplexity "5" and init-mode "random"
Generating plot with perplexity "10" and init-mode "random"
Generating plot with perplexity "20" and init-mode "random"
Generating plot with perplexity "30" and init-mode "random"
Generating plot with perplexity "40" and init-mode "random"
Generating plot with perplexity "5" and init-mode "pca"
Generating plot with perplexity "10" and init-mode "pca"
Generating plot with perplexity "20" and init-mode "pca"
Generating plot with perplexity "30" and init-mode "pca"
Generating plot with perplexity "40" and init-mode "pca"
```



0.3 Exercise 3

a)

- First, given the two points x_i & x_j the probability P_{ij} is computed to model the similarity between both points in high dimensional space.
- Second, a Student's t-Distribution with one degree of freedom is used to compute joint-probability distribution in lower dimension Q , where data positions should be optimized.
- Last, gradient of Kullback-Leibler divergence between P and Q is used to minimize cost function C .

A-tSNE introduces the approximations at initialization stage, by using approximated KNN computations to build the approximated high dimensional joint-probability distribution p_A instead of exact distribution P . The cost function C is minimized to obtain the approximated embedding described by Q_A . The similarity between points is calculated with approximated neighbours N_A , instead of exact neighborhood N_i .

b) The idea of Progressive Visual Analysis is to provide intermediate results with which the user can start the analysis process, in case the computation of the final results is costly.

Given that the distances in high-dimensional space need to be precomputed to start minimization process, the initialization process could take time before the first intermediate result is generated, and any adjustment to data will force the user to wait for reinitialization.

A-tSNE approximates the distances using KNN queries instead of precomputing them. This allows the computation of iterative minimization to start instantly and immediate processing of data.

c) As the data acquisition process could be imperfect, and the data could be incomplete or noisy and therefore data pre-processing is required.